

Names

March 28, 2020

Outline

The Problem

Naming Concepts

Name Spaces

Name Resolution

Additional Reading

Server/Object Location

Problem: How does a client know where is the server?

Solution: Not one, but several alternatives:

- ▶ *hard coded*, seldom;
- ▶ via program arguments: more flexible, but ...;
- ▶ via configuration file;
- ▶ via *broadcast/multicast*;
- ▶ via a location/name service:
 - ▶ local, e.g. *rmiregistry*.
 - ▶ *global*.

Addresses vs. Names

- ▶ Names are ... sequences of symbols (bits/characters/...) that refer to entities/objects.
- ▶ In the labs, we have used IP addresses (and ports)
- ▶ Addresses are **names** of **access points**. Or as Shoch put it:
*The **name** of a resource indicates **what** we seek,
an **address** indicates **where** it is,
(and a **route** tells us **how** to get there.)*
- ▶ Addresses have some limitations:
 - ▶ Addresses often are location dependent and change frequently
 - ▶ E.g. when a service is moved from one computer to another
- ▶ Names have some advantages over addresses:
 - ▶ They can be human-friendly.
 - ▶ They can hide both complexity and dynamics
 - ▶ E.g. they can hide access point changes
- ▶ Naming is a layer of indirection
 - ▶ Ultimately you need an address to access/operate on an object

Identifiers

- ▶ An **identifier** is a name with 3 properties:
 1. an identifier refers to one entity at most;
 2. an entity has at most one identifier;
 3. an identifier refers always to the same entity (it is never reused).
- ▶ Identifiers provide a mean to refer to an entity in a precise way, independently of its access points.
- ▶ Examples?
 - ▶ From the "real" world?
 - ▶ From the "virtual" world?

Pure Names

- ▶ Are names that contain no information whatsoever about what they refer to:
 - ▶ Not only about location, but about anything else
 - ▶ They do not commit the system to anything
 - ▶ They are useful only for comparison

Problems/challenges of pure names

- ▶ where to look them up to find out information about them?
- ▶ how do you know that an object does not exist? How can a global search be avoided?
- ▶ how to engineer uniqueness reliably in a distributed system?

Problems/challenges of impure names

- ▶ what if the information yielded by the name, e.g. location, is not valid anymore?
 - ▶ This is specially relevant for mobile systems, and requires appropriate solutions

Bindings, Contexts and Name Resolution


- ▶ A **binding** is a mapping from a name to an object/entity (usually identified by a lower-level name, e.g. address)
- ▶ A context/name space is a set of **bindings**
- ▶ A name space defines:
 - ▶ the syntax and structure (flat vs. hierarchical) of a name
 - ▶ the rules to find a binding of a name (**name resolution**)
- ▶ **Name resolution** is the process of finding a binding for a name
- ▶ A name is always resolved in the context of its name space:

file name	->	OS filesystem
Java program variable	->	JVM executing the program
ISBN of a publication	->	ISBN (Intern.Standard Book Number)
Car license plates	->	national/regional license plate regist


Name Resolution in a Distributed System

- ▶ Usually, name resolution is done with the help of a name service
- ▶ In small scale distributed systems, name resolution requires only one server:
 - ▶ E.g., the `rmiregistry`
- ▶ In distributed systems of larger scale, name resolution may require more than one server. In this case, name resolution can use one of 3 strategies:
 1. Iterative
 2. Recursive.
 3. Transitive.


Name Resolution: Strategies



Iterative



Recursive



Transitive

- ▶ Recursive name resolution:
 - ▶ Allows for caching at servers
 - ▶ This may make resolution more efficient (with lower communication costs)
 - ▶ But, it:
 - ▶ requires servers to keep state
 - ▶ makes it harder to set the values of timeouts
- ▶ Transitive name resolution also makes it harder for the client to set a timeout value

Name Resolution and *Closure Mechanism*

Names are resolved always in a context

Problem

- ▶ How do you get a context that you can use to resolve a name?
 - ▶ How do you get a "remote reference to the `rmiregistry`"?
 - ▶ How to start the name resolution of a name of a file system:
i.e. where is the root directory?
 - ▶ How to find the IP address of a DNS server to resolve a DNS name?

Response

Use a **closure mechanism**

- ▶ Typically this is an *ad-hoc* and simple solution.

Hierarchical Name Spaces

- ▶ Most name spaces have a hierarchical structure:
 - ▶ OS filesystem
 - ▶ Domain Name System (DNS)
 - ▶ Postal addresses
 - ▶ Car license plates are resolved in another context – per country, region etc.
- ▶ A hierarchical structure simplifies:
 - ▶ the assignment;
 - ▶ the resolutionof names
- ▶ Allows to partition a name space into naming domains
 - ▶ Often, a naming domain has an administrative authority for assigning names within it
 - ▶ An administrative authority may delegate name assignments for sub-domains (e.g. in DNS)

Additional Reading

- ▶ Chapter 5 of van Steen and Tanenbaum, *Distributed Systems, 3rd Ed.*
 - ▶ Section 5.1: *Names, Identifiers and Addresses*
 - ▶ Section 5.3: *Structured Naming*
- ▶ J. Saltzer, *On the Naming and Binding of Network Destinations*, in RFC 1498, 1993